

Simulation and Scientific Computing 2 Seminar

Kristina Pickl, Dominik Bartuschat, Christoph Rettinger, Ulrich Rüde
SS 2015
Chair for System Simulation (LSS)



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

Outline

Full Multi-Grid (FMG)

Introduction to Parallel Programming

Distributed vs. Shared Memory

OpenMP

MPI-Message Passing Interface

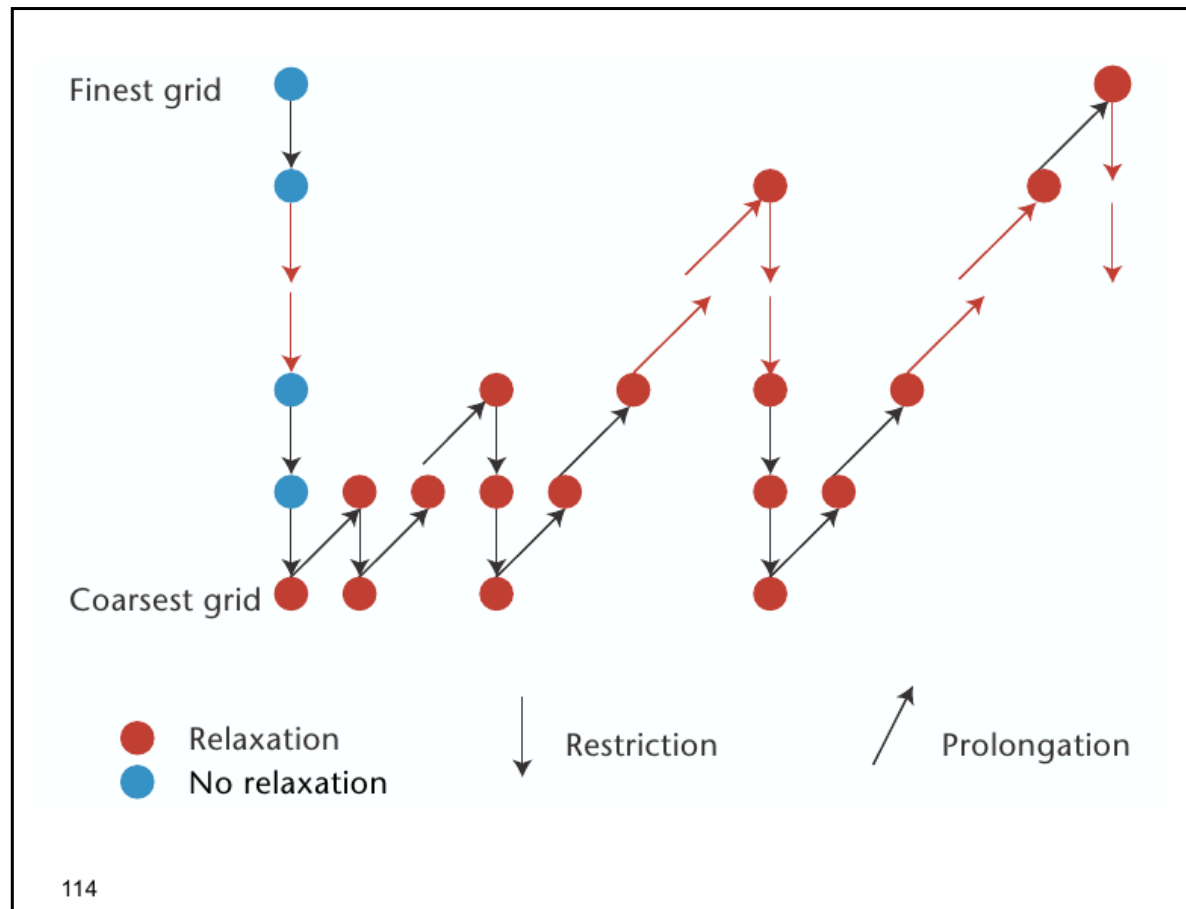
Full Multi-Grid (FMG)



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

Finite State Machine






Introduction to Parallel Programming




FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

Distributed vs. Shared Memory



- Hardware
 - same program on each processor/machine
 - explicit programming required
- Programming
 - all variables process-local, no implicit knowledge of data on other processors
 - send/receive messages of suitable library
- Languages
 - e.g. MPI
 - <http://www.mpi-forum.org/>



- Hardware
 - single program on single machine
 - workload distributed among threads
- Programming
 - all variables either shared among all threads or duplicated for each thread
 - threads communicate by sharing variables
- Languages
 - e.g. OpenMP
 - <http://www.openmp.org/>

OpenMP Programming Model

- based on the `#pragma` compiler directives

```
#pragma omp directive [clause list]
```

- OpenMP programs execute serially until they encounter `parallel` directive

```
/* serial code segment
```


```
#pragma omp parallel [clause list]
{
    /* parallel code segment
}
```

```
/* rest of serial code segment */
```

- each thread executes structured block specified by the `parallel` directive
- the `clause list` specifies the conditional parallelization, number of threads and data handling

Execution Model

Fork-Join Model



Data Handling

- `private (variable list)`: indicates that the set of variables is local to each thread (i.e. each thread has its own copy of each variable in the list)
- `shared (variable list)`: indicates that all variables in the list are shared across all the threads, i.e., there is only one copy

```
#include <omp.h>
int main()
{
    int i=5; // a shared variable

    #pragma omp parallel
    {
        int c; // a variable private to each thread
        c = omp_get_thread_num();
        std::cout << "c: " << c << ", i: " << i << std::endl;
    }
}
```

Data Handling

- `reduction (operator list)`: specifies how multiple local copies of a variable are combined into a single copy at the master when threads exit
- possible operators: `+`, `*`, `-`, `&`, `|`, `&&`, `||`

```
#include <omp.h>
int main()
{
    double sum ( 0.0 );

    #pragma omp parallel reduction (+: sum)
    {
        /* compute local sums here */
    }
    /* sum here contains sum of all local instances of sum */
}
```

Synchronization Constructs

- `#pragma omp barrier`
 - synchronizes all threads in the team
- `#pragma omp single`
 - only executed by exactly one thread
 - all other threads skip this region
 - implicit barrier at end of single construct
- `#pragma omp master`
 - only executed by the master thread
 - all other threads skip this region
 - no implicit barrier associated
- `#pragma omp critical`
 - executed by all threads
 - but only one thread at a time

Runtime Library Functions

- Setting total number of threads

- at runtime: via `omp_set_num_threads`

```
#include <omp.h>
void omp_set_num_threads( int num_threads )
```

- via environment variable `OMP_NUM_THREADS`

```
export OMP_NUM_THREADS=4
```

- getting total number of threads

```
#include <omp.h>
void omp_get_num_threads( void )
```

- getting ID of specific thread

```
#include <omp.h>
void omp_get_thread_num( void )
```

MPI Programming Model

- each processor runs a sub-program
- variables of each sub-program have
 - the same name
 - but different locations and different data
 - i.e. all variables are private
- communicate via special send and receive routines

Hello World

```
#include <mpi.h>
int main( int argc, char **argv )
{
    // Definition of the variables
    int size; //The total number of processes
    int rank; //The rank/number of this process

    // MPI initialization
    MPI_Init( &argc, &argv );

    // Determining the number of CPUs and the rank for each CPU
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    ...
}
```

Hello World

```
...  
// 'Hello World' output for CPU 0  
if( rank == 0 )  
    std::cout << "Hello World" << std::endl;  
  
// Output of the own rank and size of each CPU  
std::cout << "I am CPU " << rank << " of " << size << " CPUs" << std::endl;  
  
// MPI finalizations  
MPI_Finalize();  
return 0;  
}
```

Blocking Operations

- some operations may block until another process acts:
 - synchronous send operation blocks until receive is posted
 - receive operation blocks until message is sent
- send buffer may be reused after MPI_Send returns
- MPI call returns after completion of the corresponding send/receive operation

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest,
             int tag, MPI_Comm comm)
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source,
             int tag, MPI_Comm comm, MPI_Status *status)
```


Non-Blocking Operations

- return immediately and allow the sub-program to perform other work
- at some later time the sub-program must test or wait for the completion of non-blocking operation
- all non-blocking operations must have matching wait (or test) operations
- a non-blocking operation immediately followed by a matching wait is equivalent to a blocking operation

```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest,  
             int tag, MPI_Comm comm, MPI_Request *request)  
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source,  
             int tag, MPI_Comm comm, MPI_Request *request)
```

Collective Communications

- Broadcast operation: a one-to-many communication

```
int MPI_Bcast( void *buffer, int count, MPI_Datatype datatype,
               int root, MPI_Comm comm )
```

- Reduction operation: combine data from several processes to produce single result

```
int MPI_Reduce( void *sendbuf, void *recvbuf, int count,
                MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm )
```

- Barriers: synchronizes processes, blocks until all processes in the communicator have reached this routine

```
int MPI_Barrier( MPI_Comm comm )
```

- Wait: waits for an MPI send or receive to complete

```
int MPI_Wait ( MPI_Request *request, MPI_Status *status)
```



Thank you for your attention!



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

References

- MPI:

- <http://www.mpi-forum.org/>
- <https://www10.informatik.uni-erlangen.de/Teaching/Courses/WS2014/SiWiR/exerciseSheets/ex03/mpi.pdf>

- OpenMP:

- <http://www.openmp.org/>
- <https://computing.llnl.gov/tutorials/openMP/>
- https://www10.informatik.uni-erlangen.de/Teaching/Courses/WS2014/SiWiR/exerciseSheets/ex02/ex02_1.pdf